

CSC 374/407: Computer Systems II

Lecture 3
Joseph Phillips
De Paul University

2014 July 19

Copyright © 2011 Joseph Phillips
All rights reserved

Reading

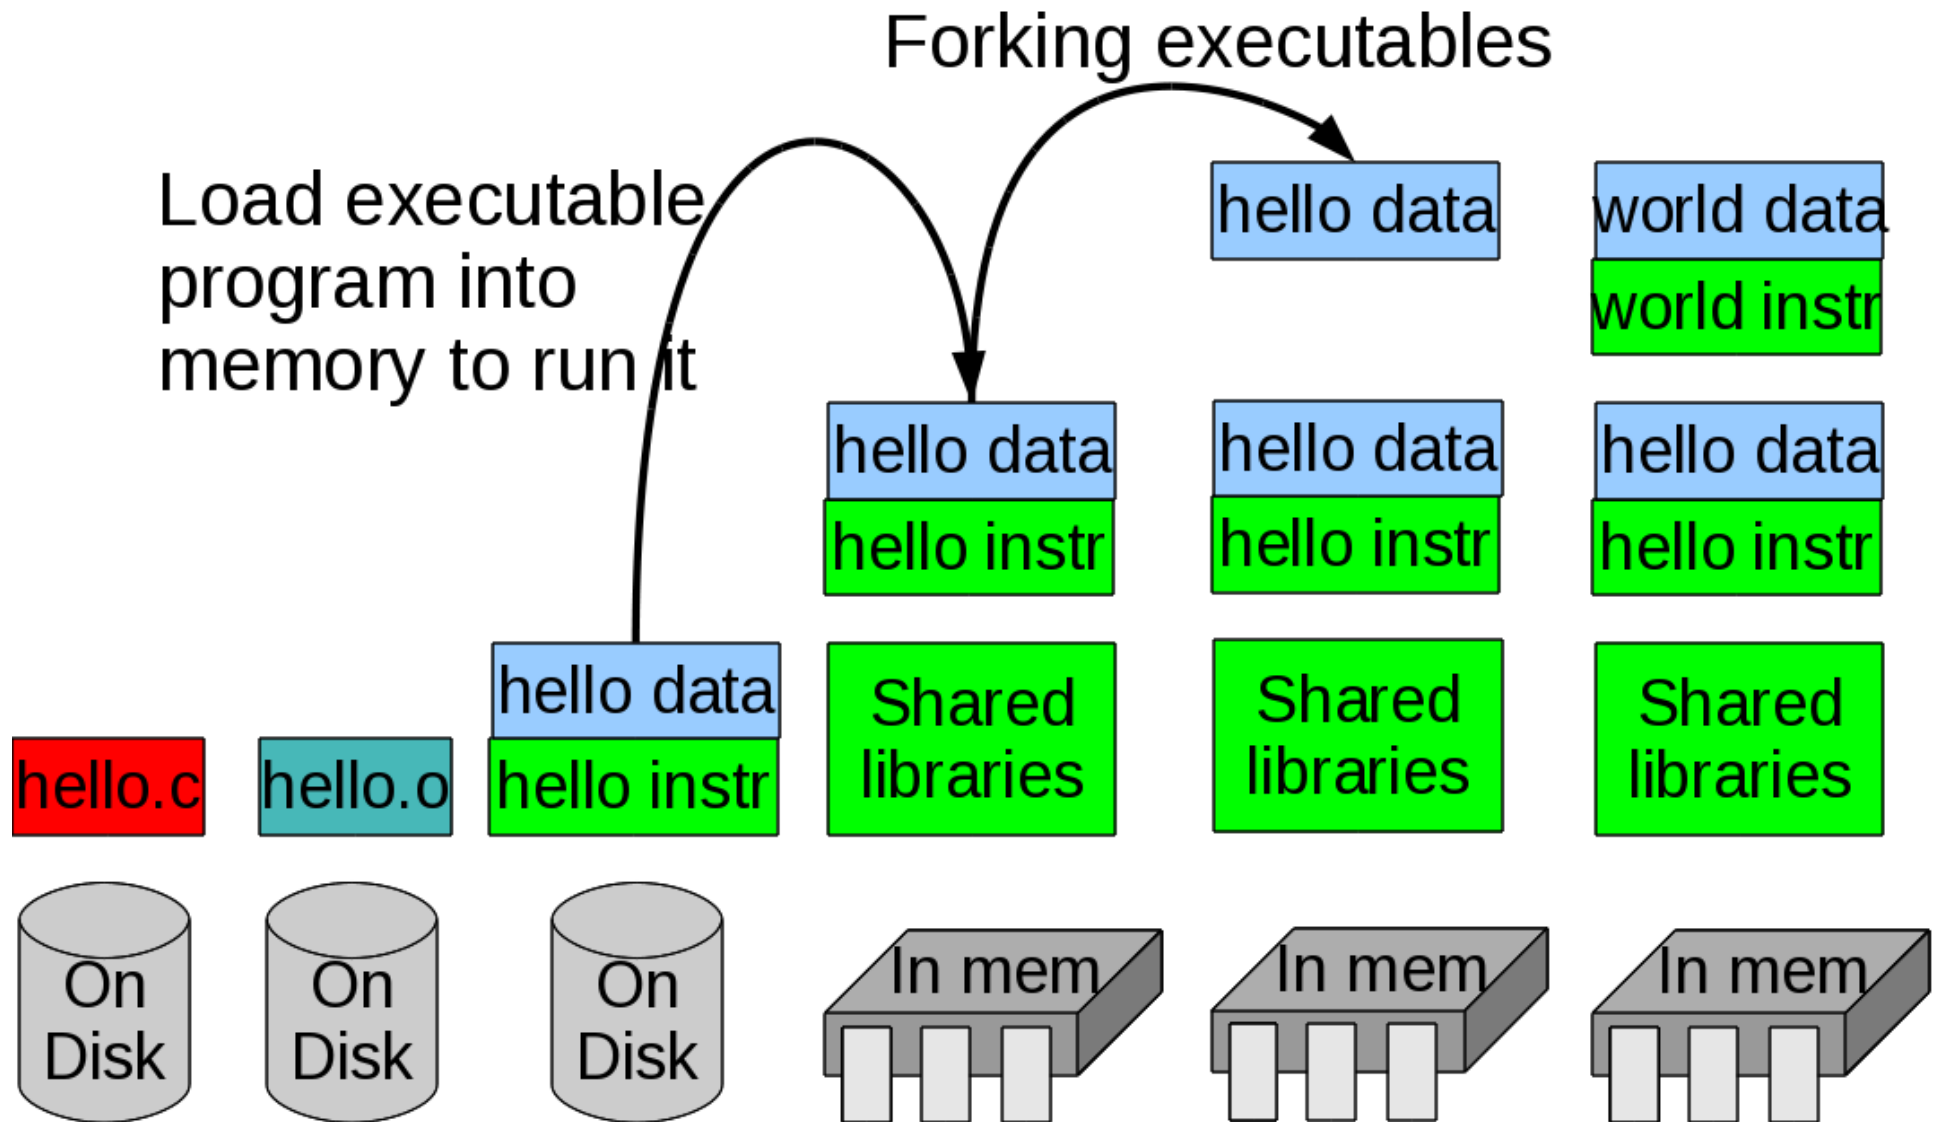
- ♦ Bryant & O'Hallaron “*Computer Systems, 2nd Ed.*”
 - ♦ Chapter 8: Exception Control Flow
- ♦ Hoover “*System Programming*”
 - ♦ System Calls 7.1-7.4

Topics

- ◆ Processes 1
- ◆ Exceptions
 - ◆ Interrupts
 - ◆ Trap
 - ◆ Fault
 - ◆ Abort
- ◆ Processes 2
 - ◆ Process lifecycle
 - ◆ `fork()` and `getpid()`
 - ◆ `exit()` and `atexit()`
 - ◆ `execvp()`
 - ◆ `wait()`, `waitpid()`
 - ◆ Zombies

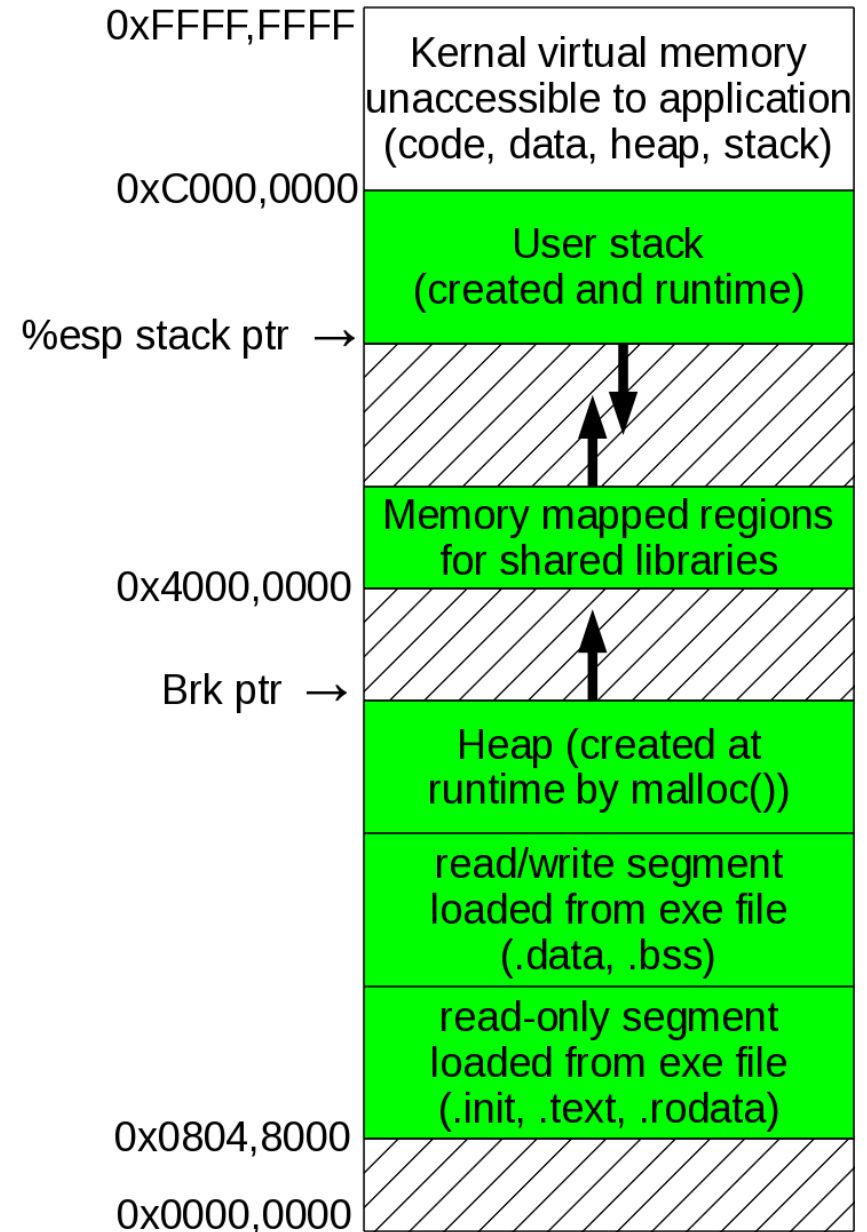
Today's topic (in time)

Running executable files and forking




Today's topic (in space)

Loading or setting up
everything:




One user's view of a computer

- Computer spends all of its time doing instructions for *my* program



Of course computers run multiple processes

- More than one user
- More than one process per user
- OS also handles hardware events (key press, network traffic, timer) and its own maintenance (disk defragmenting)



How computers do it

Load process 1

0x100: mov
0x104: inc
0x108: cmp

Save process 1

Load process 2

0x50C: push
0x510: mov
0x514: call

Save process 2

Load process 3

0xD44: pop
0xD48: add
0xD4C: ret

My stuff is
getting done!



My stuff is
getting done!



My stuff is
getting done!



What's in a context switch?

- ♦ What's has to be loaded and saved?
 - ♦ Registers
 - ♦ Page table
 - ♦ Signal vector (more about that next week)
 - ♦ CPU and memory usage stats
 - ♦ List of open files
 - ♦ *etc.*
- ♦ Doing a `jmp` to your own address or a `call` to your own function handled by a process without outside help.
- ♦ **Exceptions**, however needed to:
 - ♦ Switch between processes
 - ♦ Access system resources

4 types of exceptions

Class	Cause	Synchr onicity	Return behavior
Interrupt	Signal from hardware	Async	Next instruct
Trap	Intentional call	Sync	Next instruct
Fault	Potentially recoverable error	Sync	<i>Maybe</i> next instruct
Abort	Unrecoverable error	Sync	Never returns

Interrupts

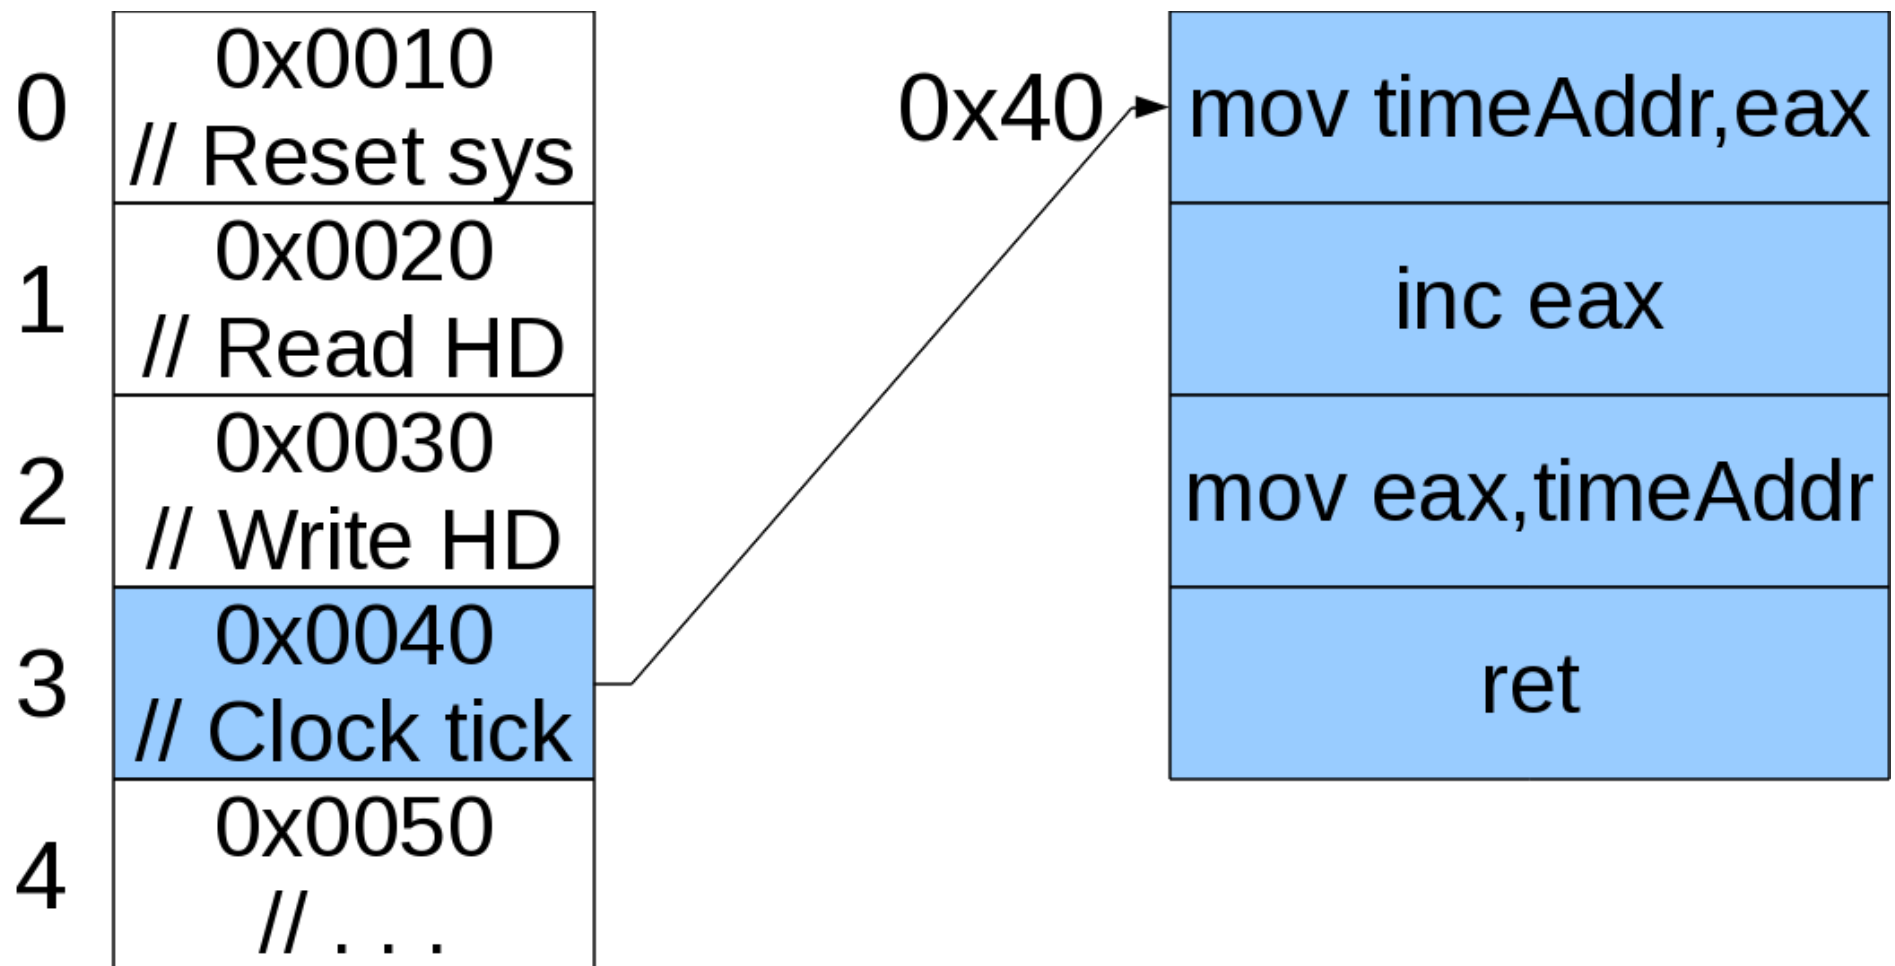
- ◆ Hardware says “I don't care which ordinary user's process you're running, ***attend to me now! (or at least soon)***”
 - ◆ A key is pressed by the user
 - ◆ A packet comes over the network
 - ◆ Computer's timer goes off

Doing an interrupt

1. An interrupt (identified by an integer) fires
2. Save state of running process
 - Common to push registers on stack
3. Go to interrupt table (*aka* vector table): array of pointers to code to handle interrupts
4. Get address from array
5. Go to that address
6. Do code
7. Re-load state of running process

Interrupt table or vector table

- Array of pointers to functions to handle specific interrupts



“IBM PC” i386 Interrupt Table (1)

- ♦ ROM BIOS by IBM (Microsoft?) circa 1981
- ♦ Now standardized and written by:
 - ♦ American Megatrends Inc
 - ♦ Micro Firmware
 - ♦ Phoenix Technologies

00h CPU: Executed after an attempt to divide by zero or when the quotient does not fit in the destination

01h CPU: Executed after every instruction while the trace flag is set

02h CPU: NMI, used e.g. by POST for memory errors

03h CPU: The lowest non-reserved interrupt, it is used exclusively for debugging, and the INT 03 handler is always implemented by a debugging program

“IBM PC” i386 Interrupt Table (2)

- 04h** CPU: Numeric Overflow. Usually caused by the INTO instruction when the overflow flag is set.
- 05h** Executed when Shift-Print screen is pressed, as well as when the BOUND instruction detects a bound failure.
- 06h** CPU: Called when the Undefined Opcode (invalid instruction) exception occurs. Usually installed by the operating system.
- 07h** CPU: Called when an attempt was made to execute a floating-point instruction and no numeric coprocessor was available.
- 08h** IRQ0: Implemented by the system timing component; called 18.2 times per second (once every 55 ms) by the PIC.
- 09h** IRQ1: Called after every key press and release (as well as during the time when a key is being held)

“IBM PC” i386 Interrupt Table (3)

0Ah Reserved for OS?

0Bh IRQ3: Called by serial ports 2 and 4 (COM2/4) when in need of attention

0Ch IRQ4: Called by serial ports 1 and 3 (COM1/3) when in need of attention

0Dh IRQ5: Called by hard disk controller (PC/XT) or 2nd parallel port LPT2 (AT) when in need of attention

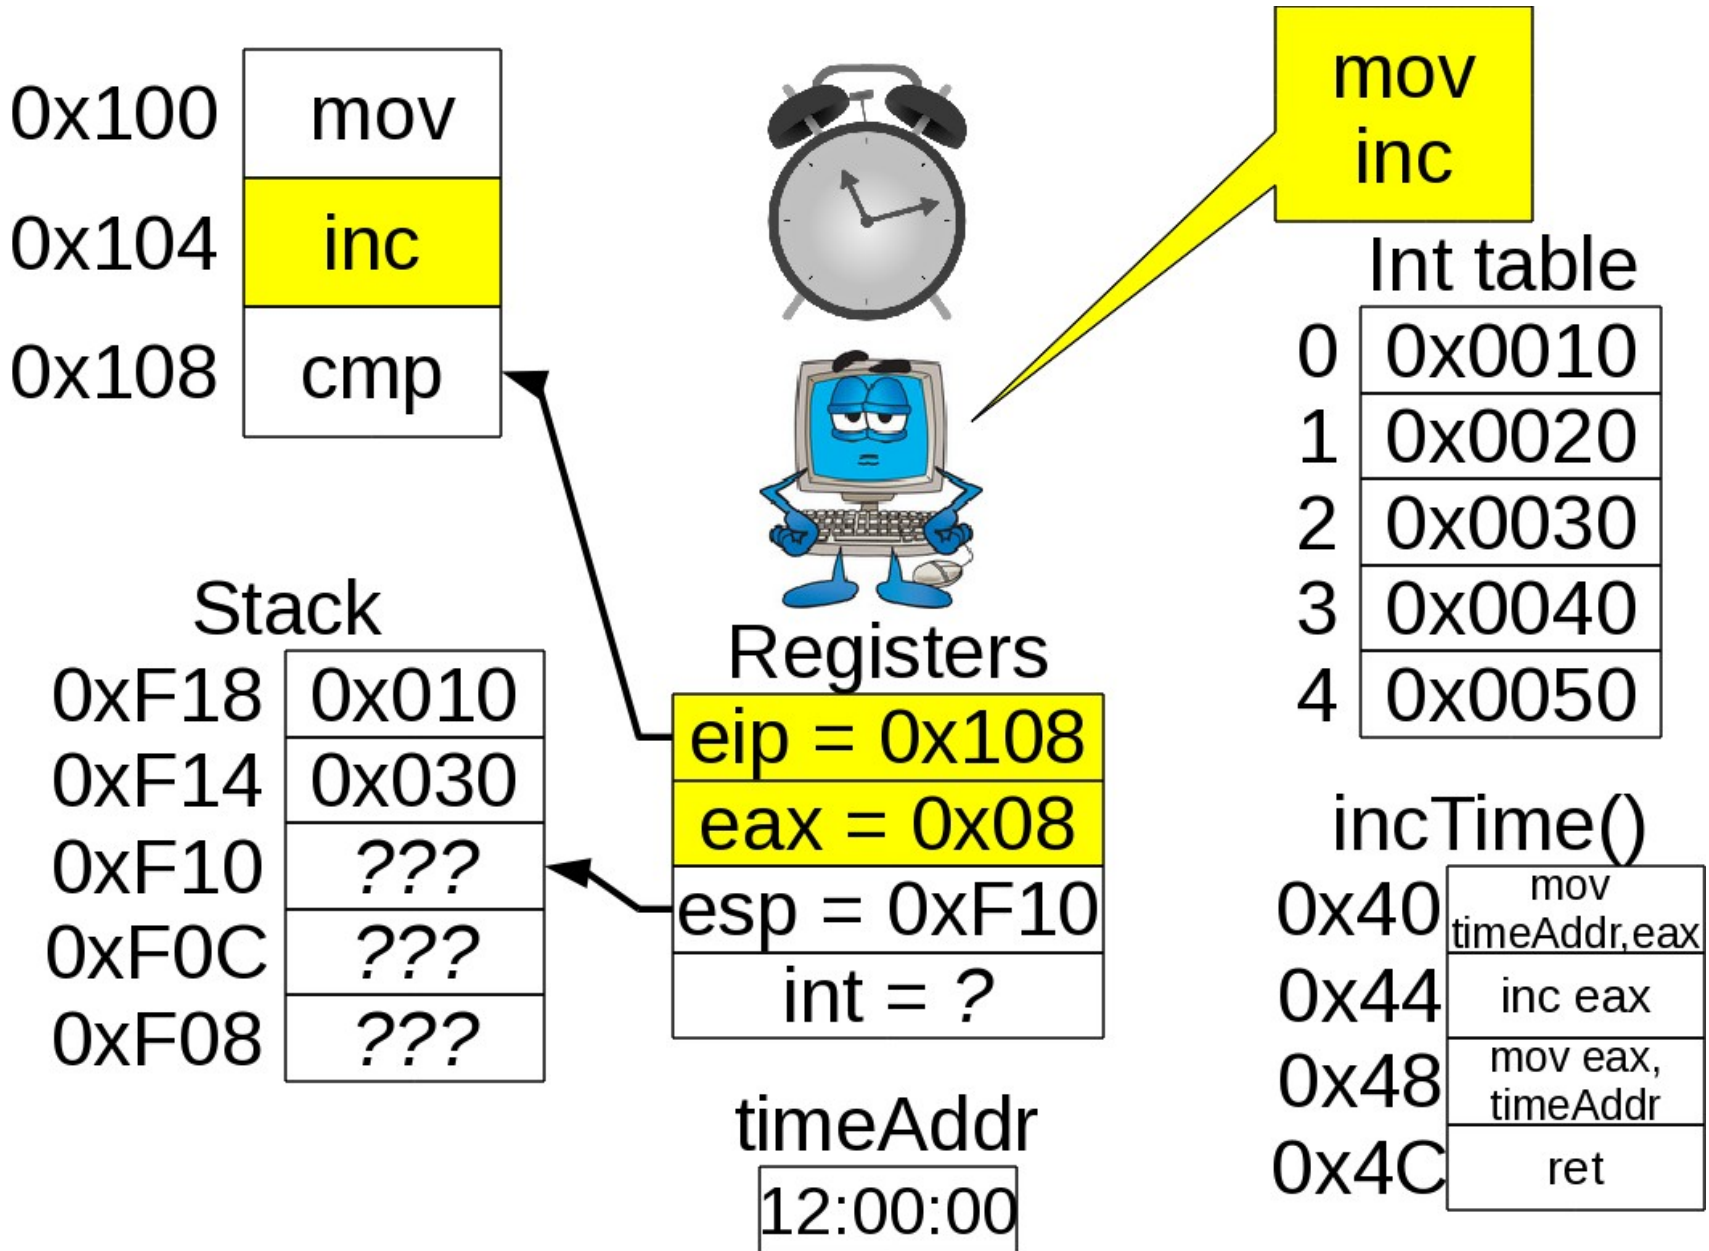
0Eh IRQ6: Called by floppy disk controller when in need of attention

0Fh IRQ7: Called by 1st parallel port LPT1 (printer) when in need of attention

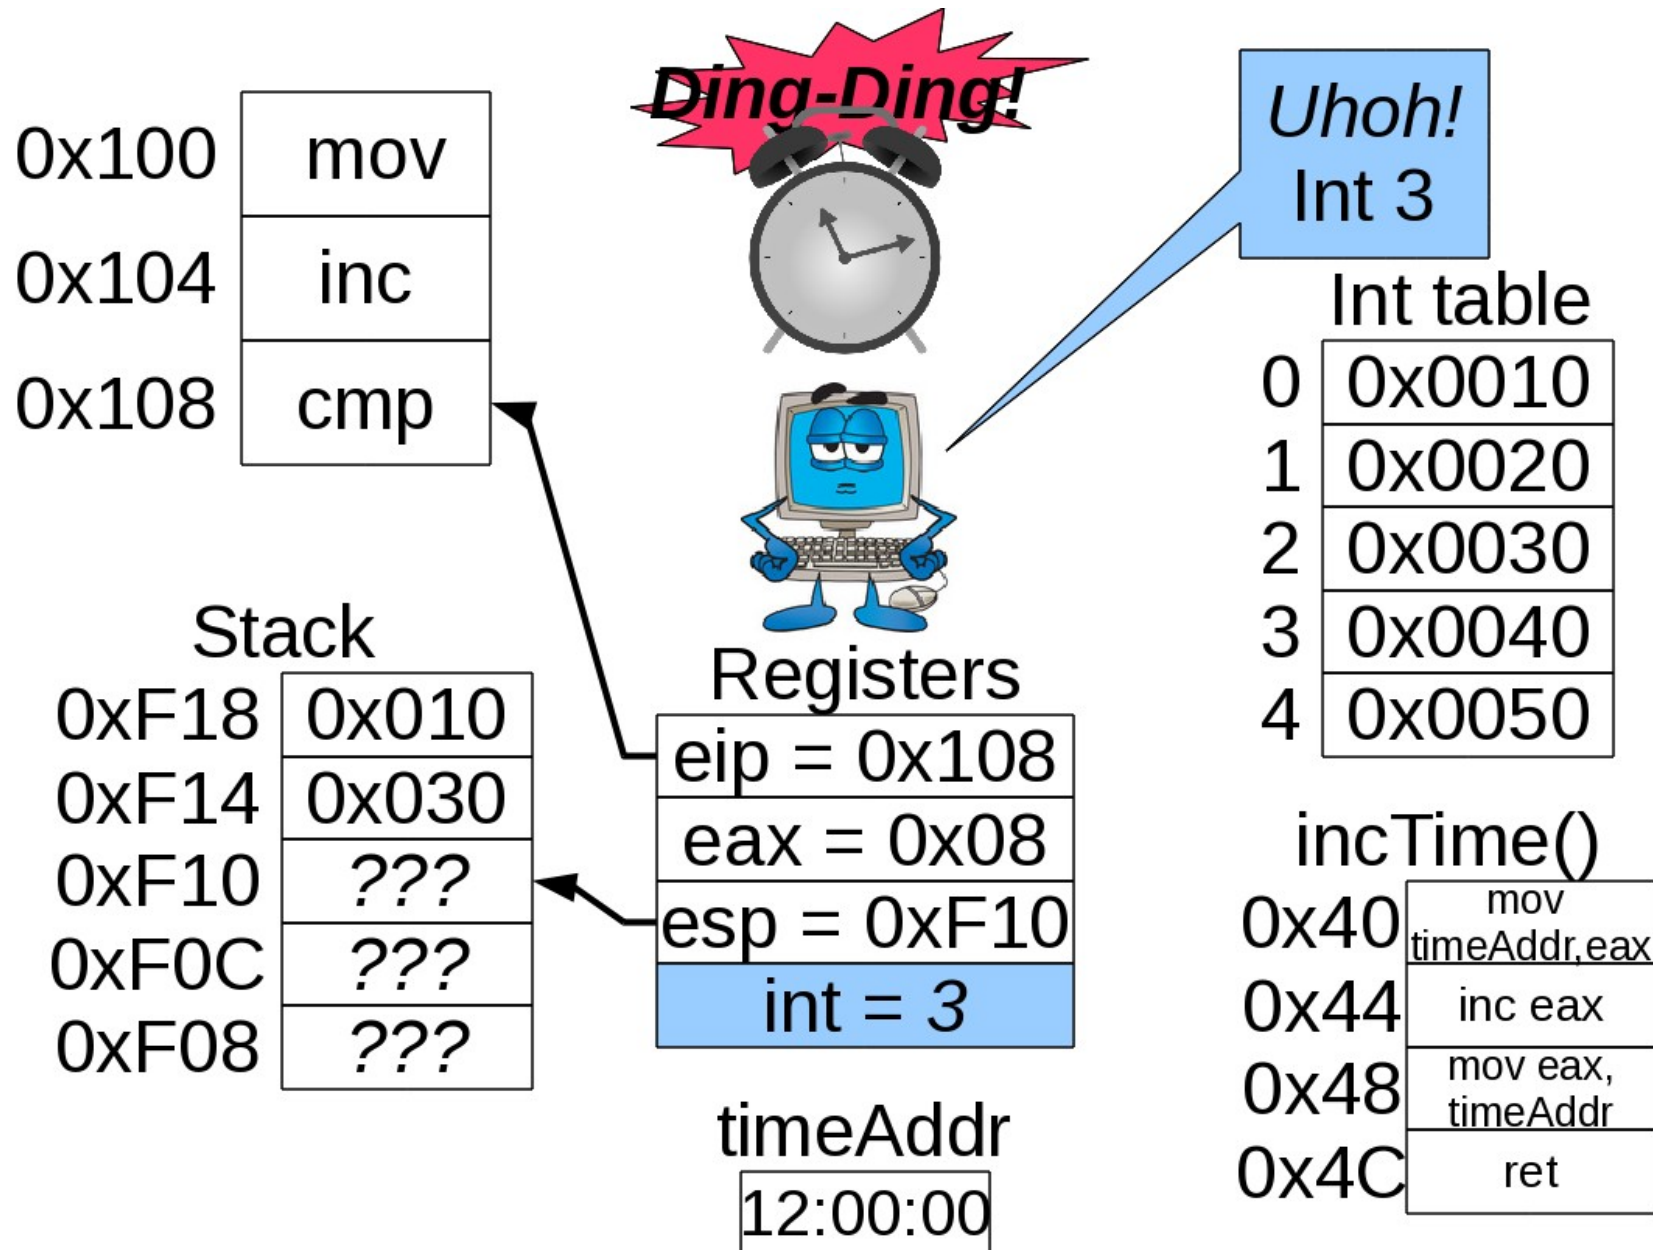
10h Video Services - installed by the BIOS or operating system; called by software programs
AH=00h: Set video mode;
AH=01h: Set Cursor shape

Etc.

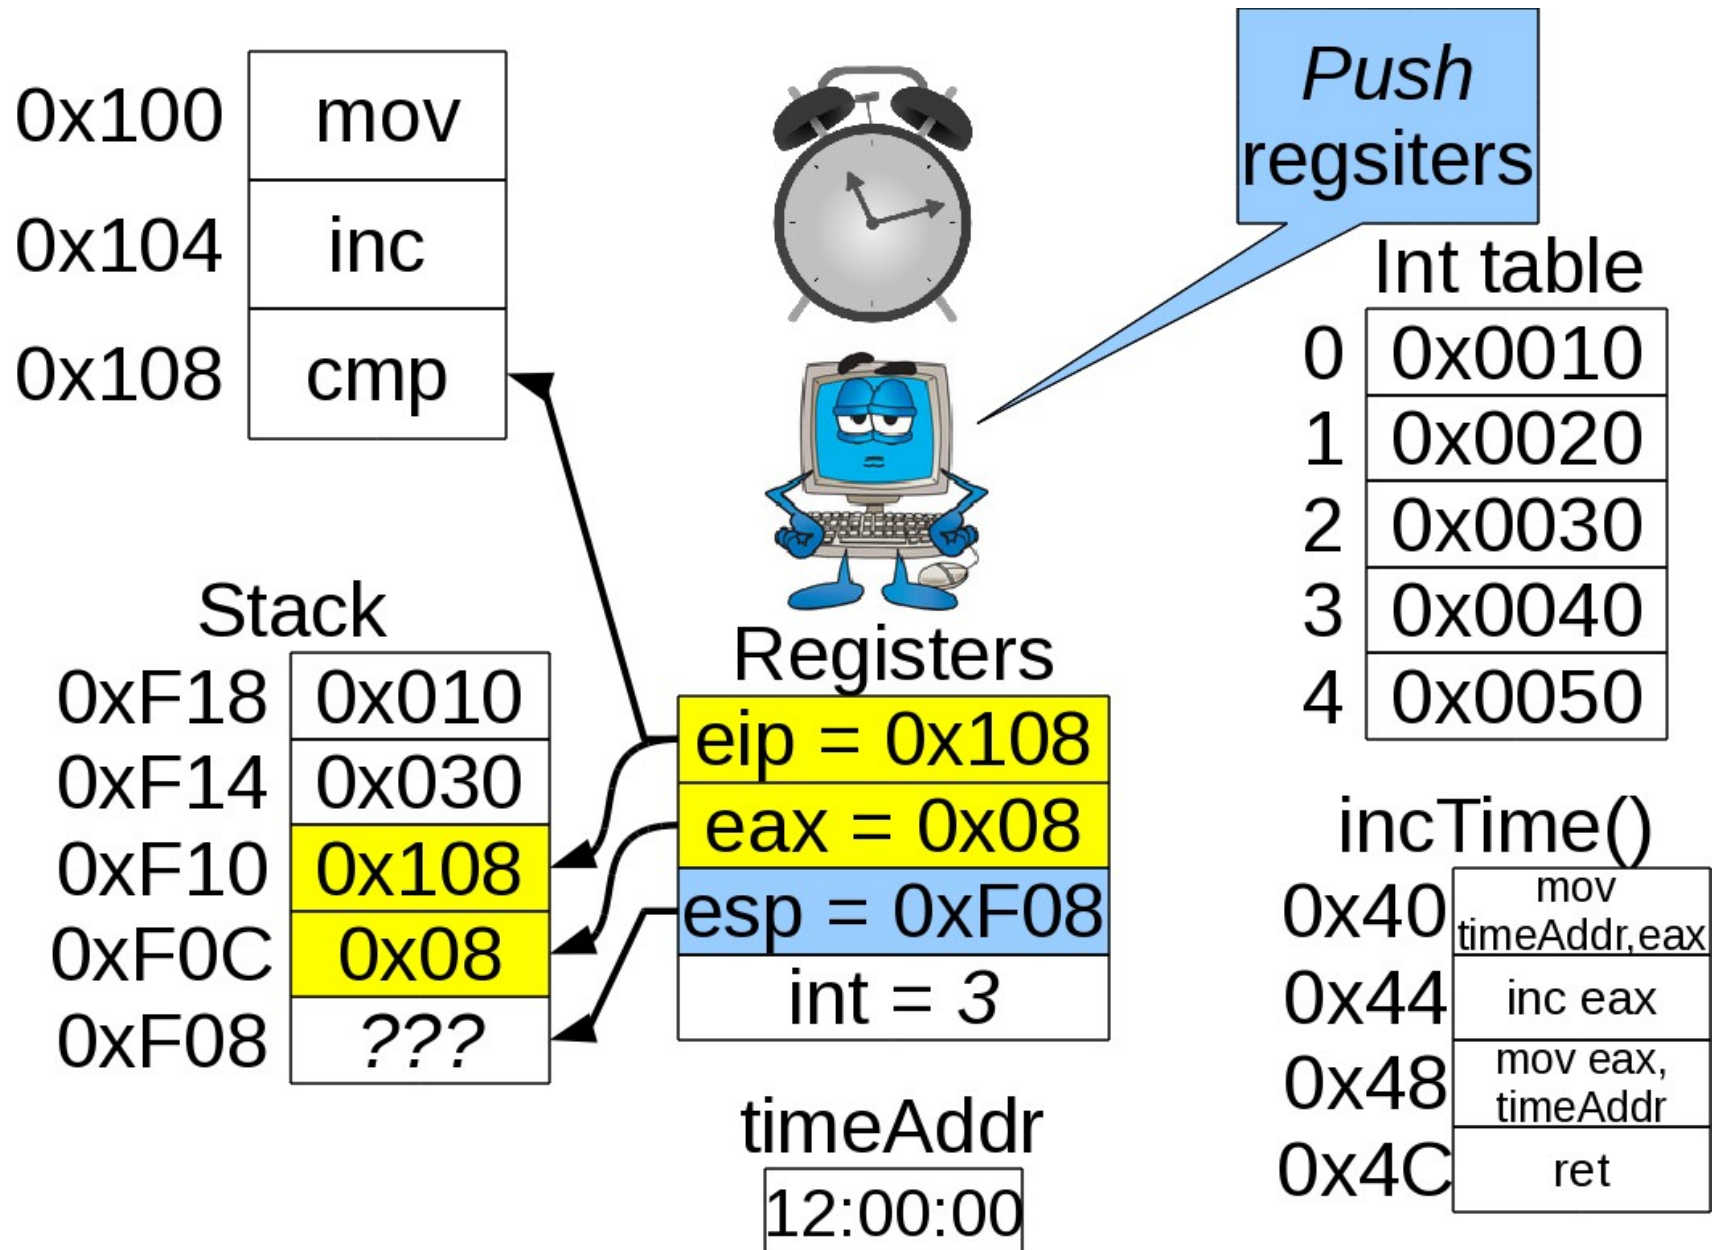
Example interrupt (1)



Example interrupt (2)



Example interrupt (3)



Example interrupt (4)

0x100	mov
0x104	inc
0x108	cmp

Stack

0xF18	0x010
0xF14	0x030
0xF10	0x108
0xF0C	0x08
0xF08	???



Get new addr
from vec table

Int table

0	0x0010
1	0x0020
2	0x0030
3	0x0040
4	0x0050

Registers

eip = 0x040
eax = 0x08
esp = 0xF08
int = 3

timeAddr
12:00:00

incTime()

0x40	mov timeAddr, eax
0x44	inc eax
0x48	mov eax, timeAddr
0x4C	ret

(b)


(c)

(a)

Example interrupt (5)

0x100	mov
0x104	inc
0x108	cmp

Stack	
0xF18	0x010
0xF14	0x030
0xF10	0x108
0xF0C	0x08
0xF08	???



Load time
Inc time
Save time

Int table	
0	0x0010
1	0x0020
2	0x0030
3	0x0040
4	0x0050

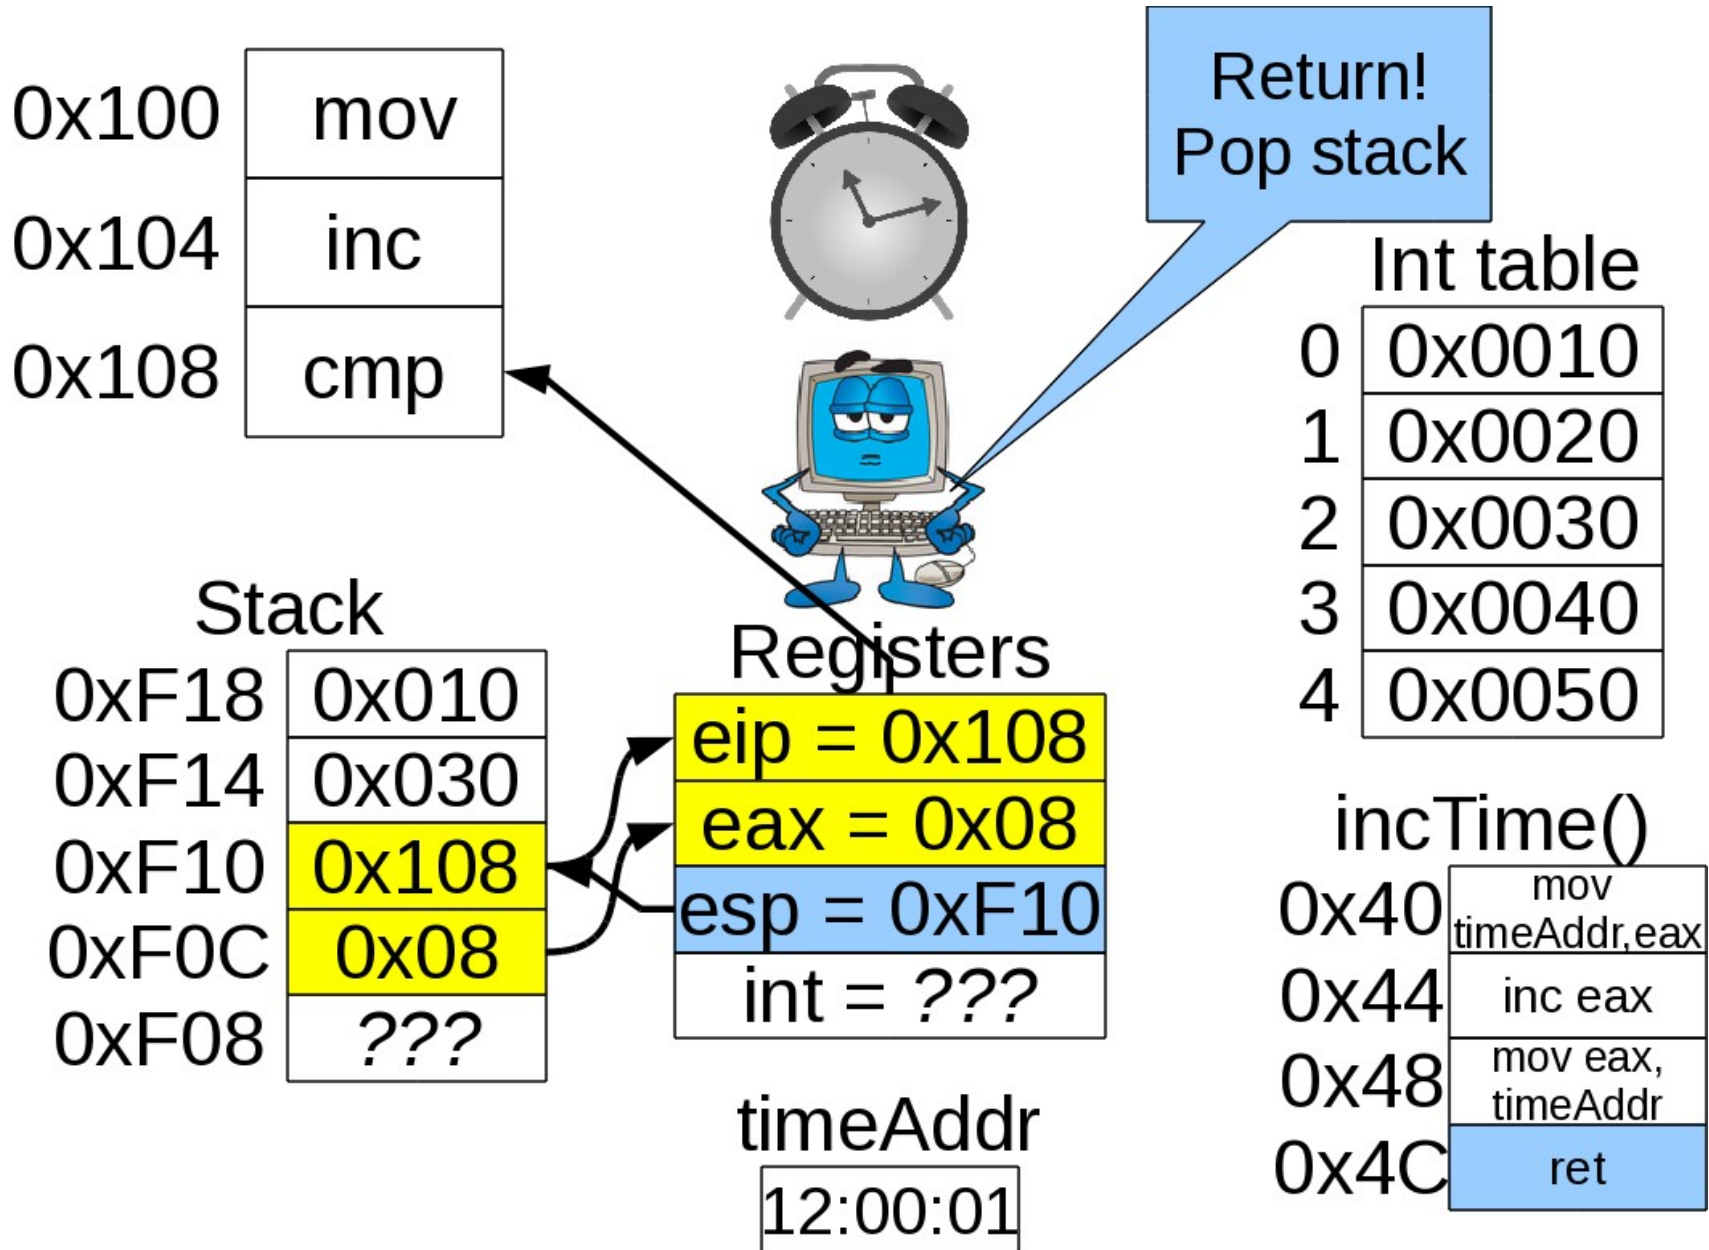
Registers	
eip	= 0x04C
eax	= 12:00:01
esp	= 0xF08
int	= 3

timeAddr
12:00:01

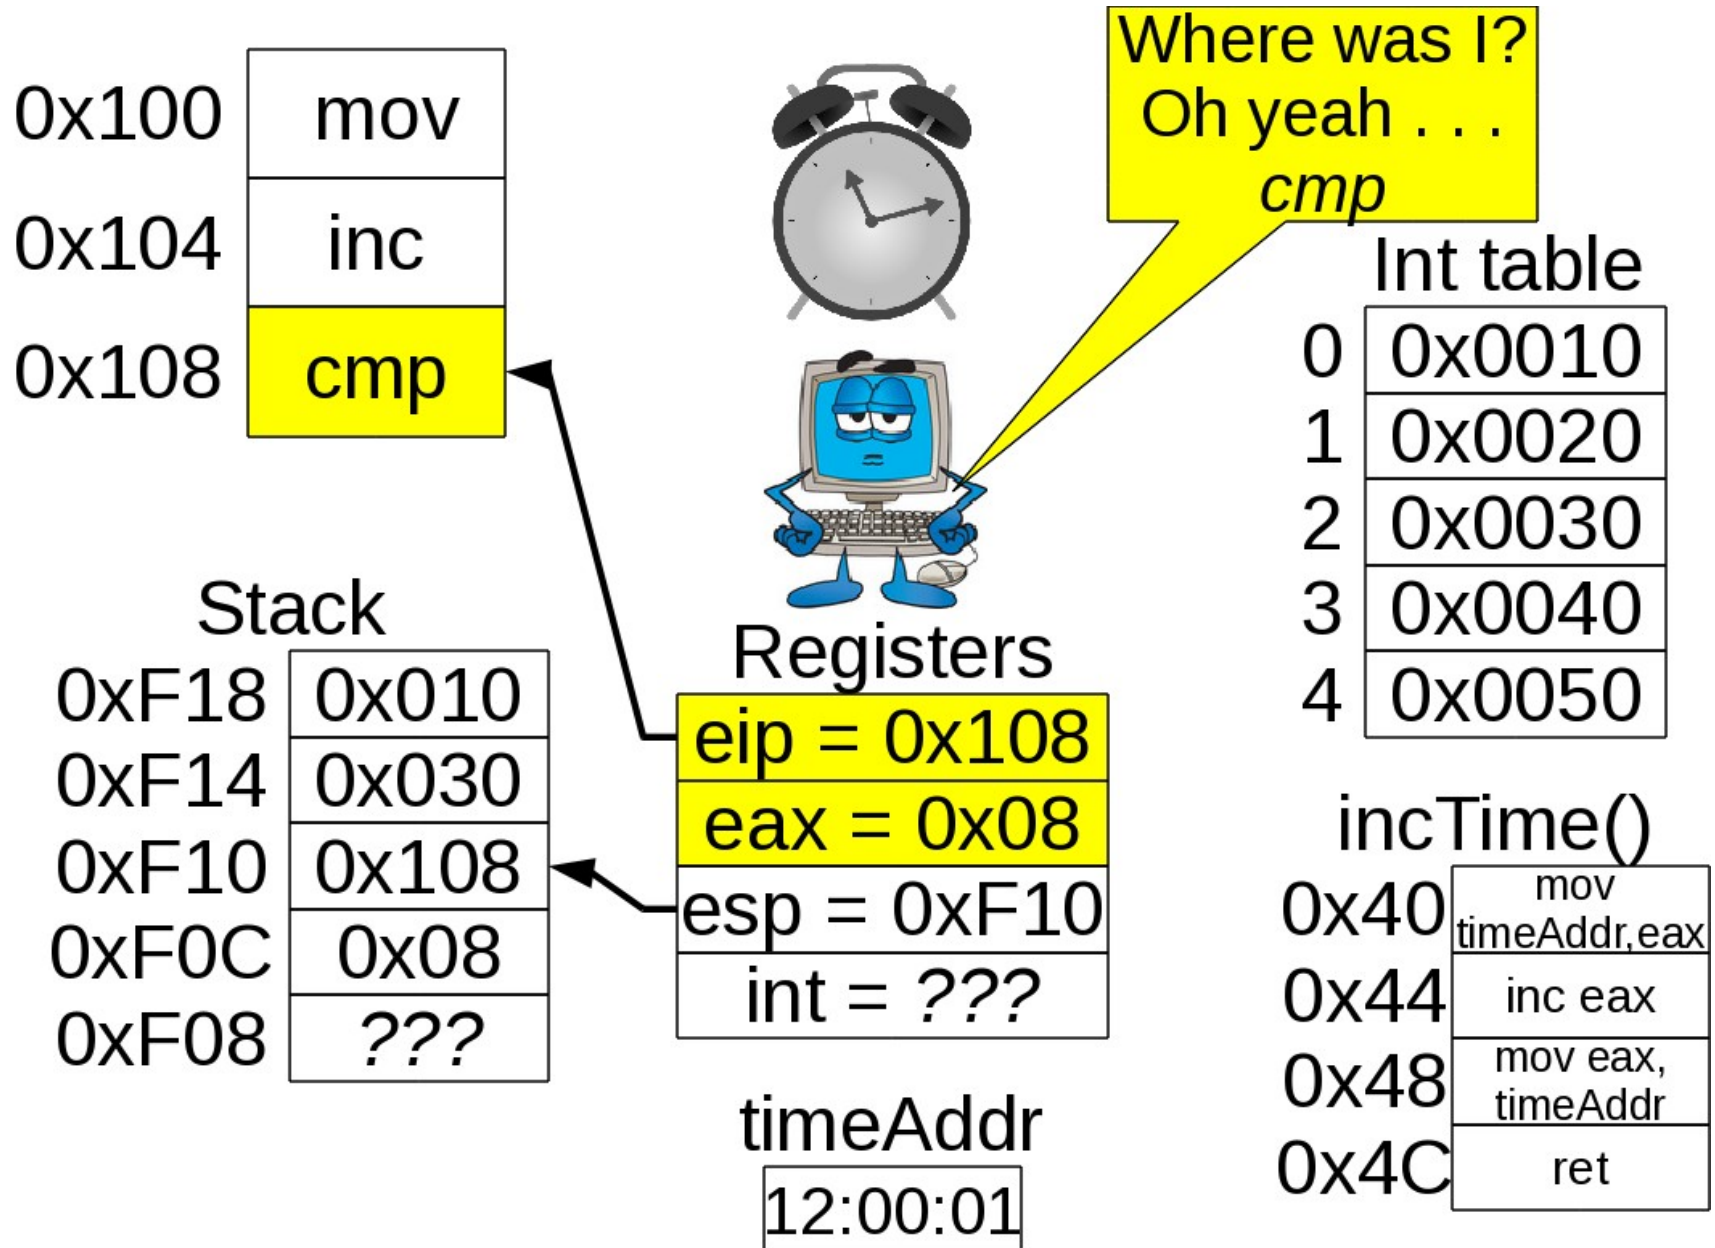
incTime()

0x40	mov
0x44	timeAddr, eax
0x48	inc eax
0x4C	mov eax,
	timeAddr
	ret

Example interrupt (6)



Example interrupt (7)



Your turn!

- ◆ The vector table on the original IBM PC was writable by any program.
- ◆ It also has a timer.
- ◆ You want to take some action every N seconds
- ◆ **Question 1:** How would you do so without messing up system time?
- ◆ **Question 2:** Is this an elegant solution?

Traps

- ♦ Intentional exceptions (e.g. system calls)
 - ♦ Linux uses `int 0x80` for all system calls

```
$ gcc -static hiWorld.c -o hiWorld
$ objdump -d -j .text hiWorld > textSeg.asm
$ grep "int " textSeg.asm
```

- ♦ Open up `textSeg.asm` and see interrupts:

```
8049109: 89 f1                mov     %esi,%ecx
804910b: 89 fa                mov     %edi,%edx
804910d: b8 92 00 00 00      mov     $0x92,%eax
8049112: cd 80               int     $0x80
```

- ♦ **Question:** If Linux always uses `int 0x80` how does it know what service to perform?
 - ♦ **Hint:** Look carefully at its preceding code

Traps, cont'd

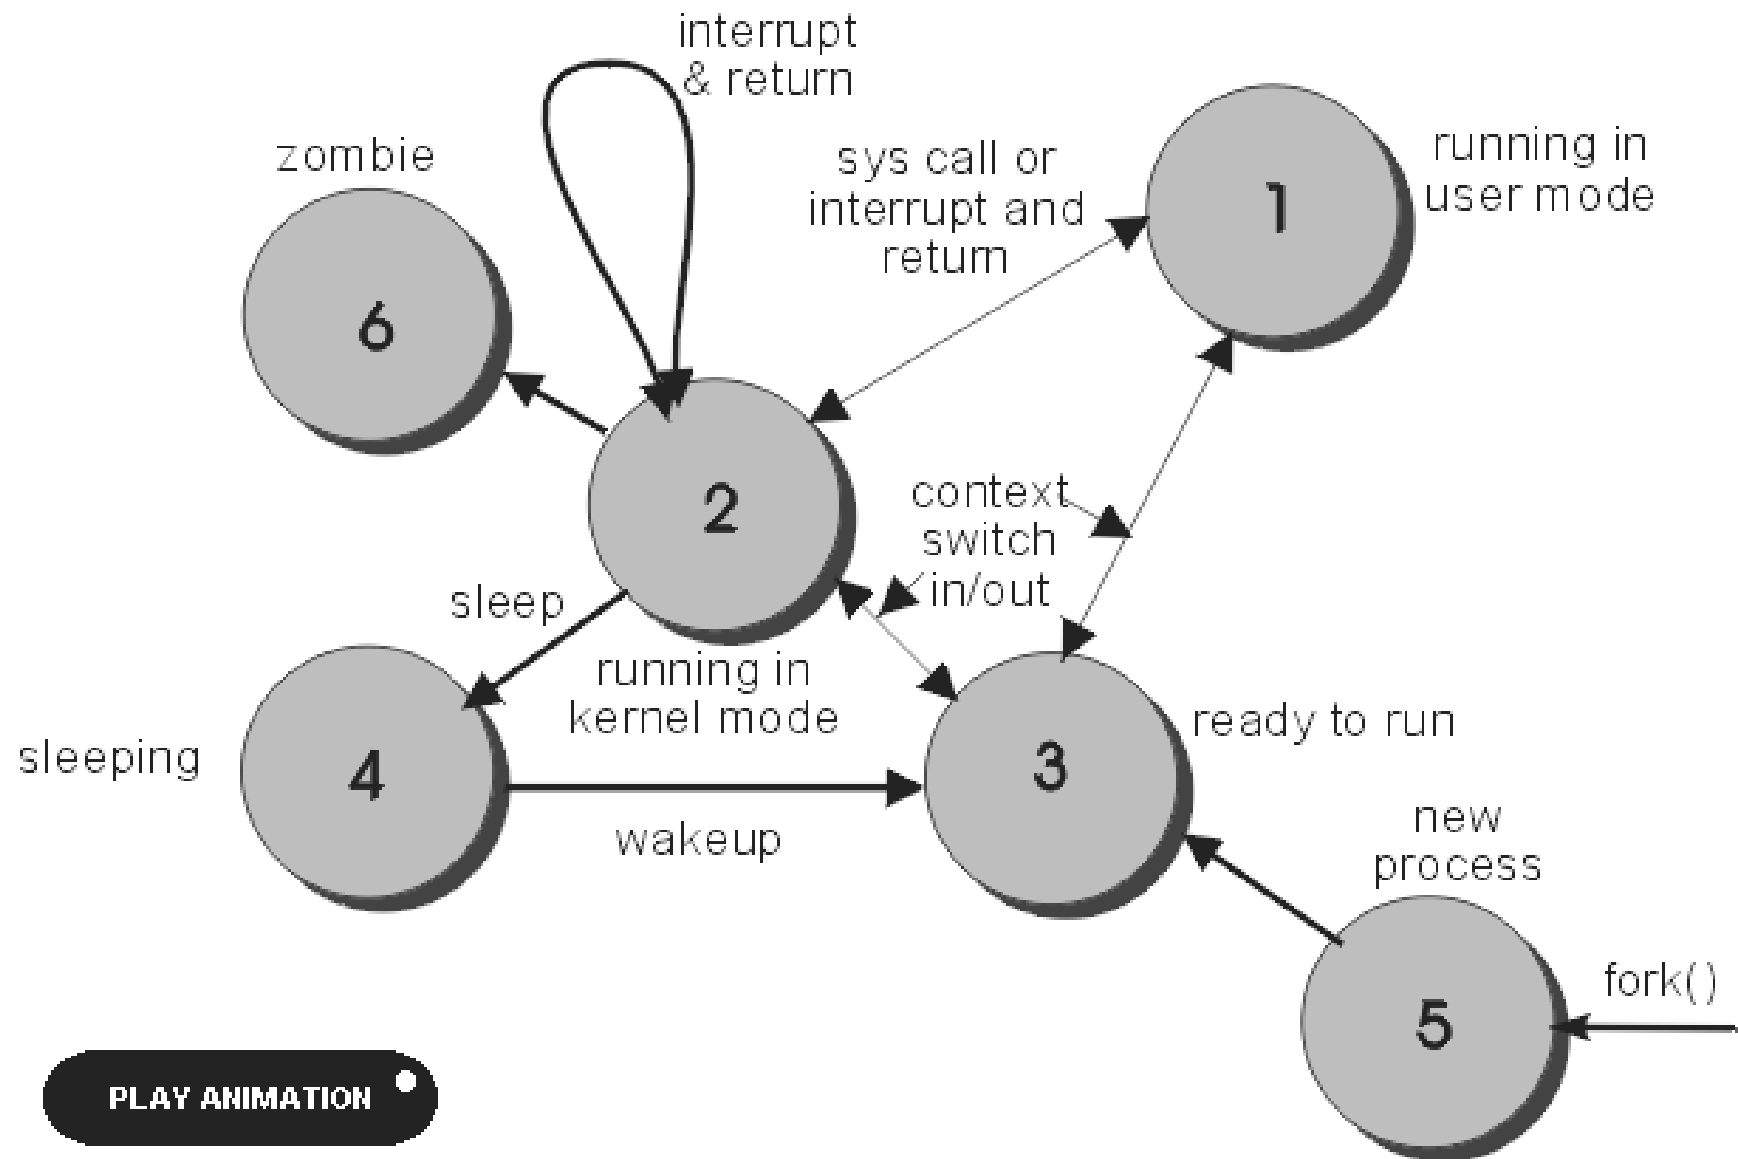
- ◆ Look in:

- ◆ `/usr/include/sys/syscall.h`, redirected to
- ◆ `/usr/include/asm/unistd.h`, finally to
- ◆ `/usr/include/asm/unistd_32.h`

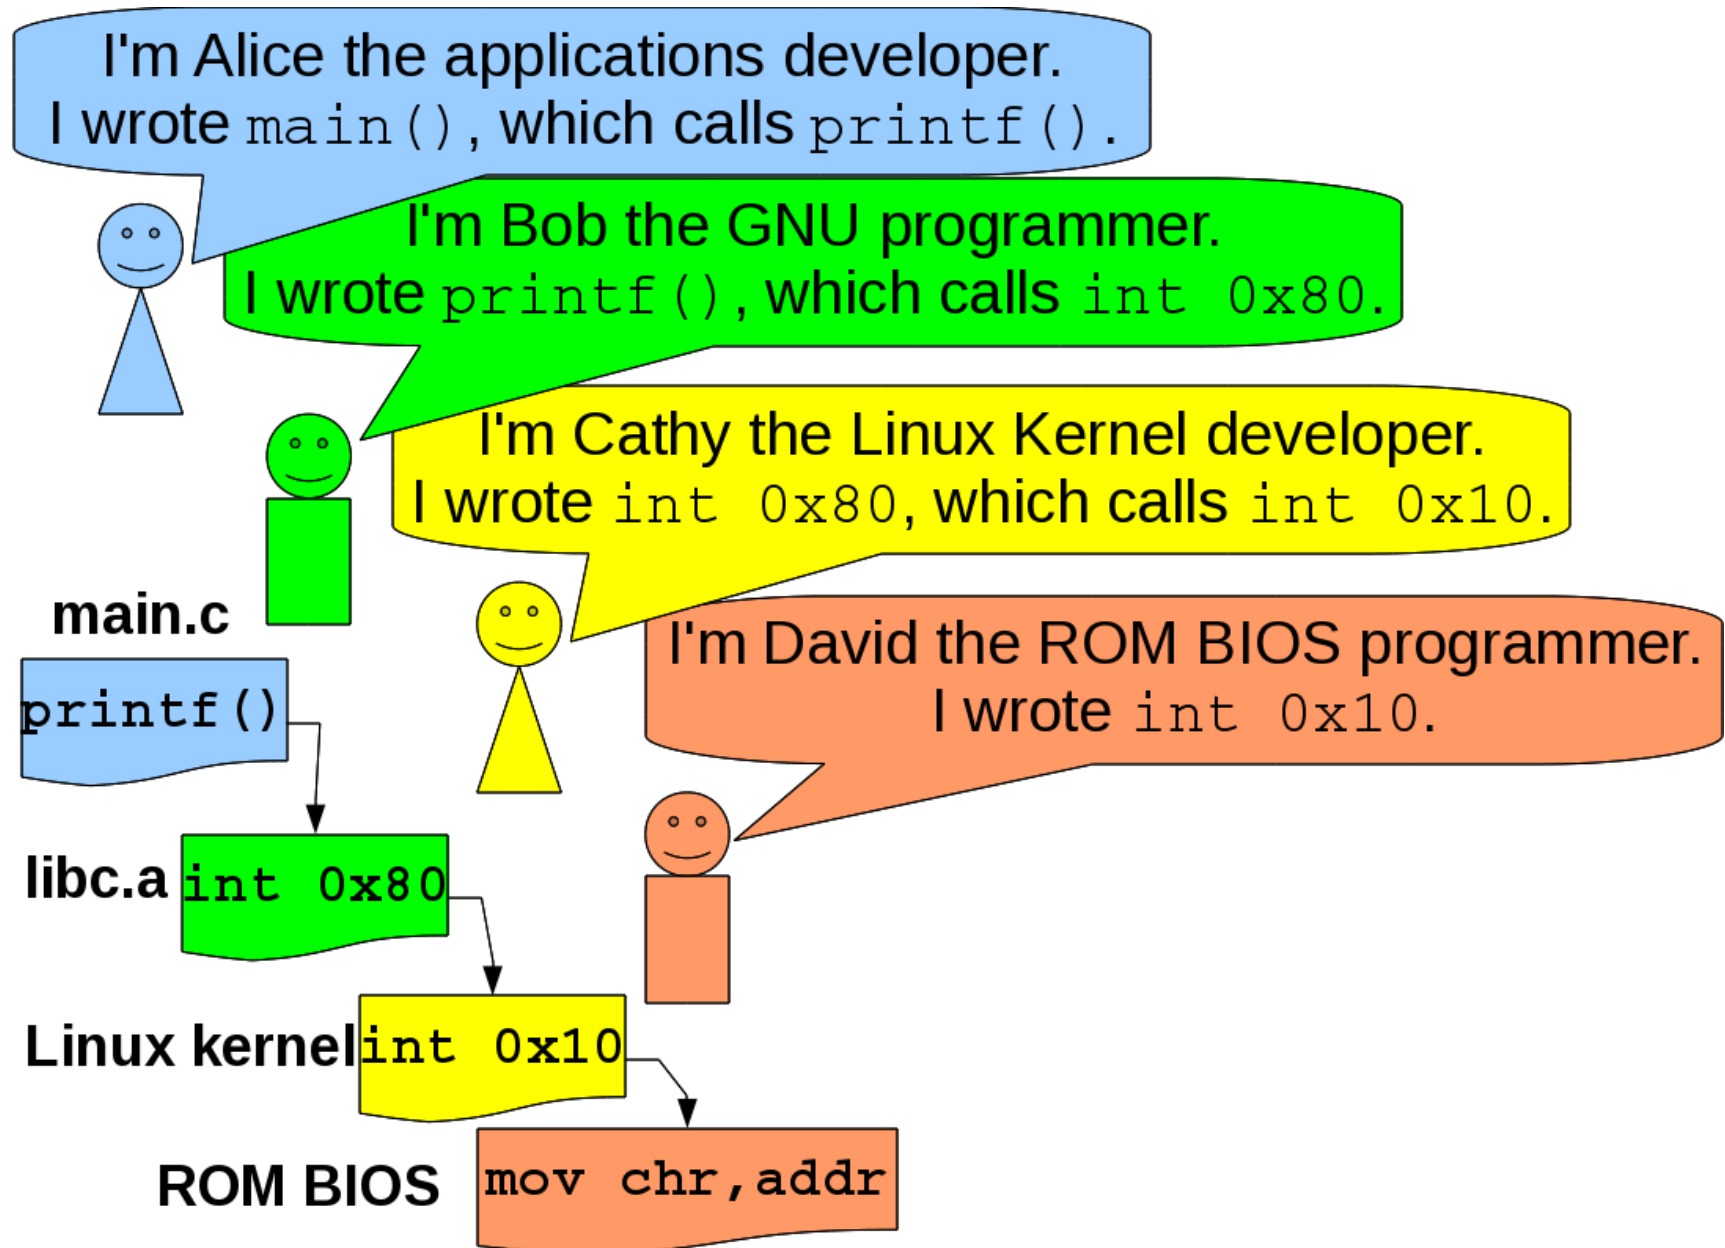
- ◆ There you'll see:

```
#define __NR_restart_syscall    0
#define __NR_exit                1
#define __NR_fork                2
#define __NR_read                3
#define __NR_write               4
#define __NR_open                5
#define __NR_close               6
#define __NR_waitpid             7
#define __NR_creat               8
. . . .
```

Doing a trap: same process in kernel mode



Following a system call



YOUR TURN!

This is such a good idea that Microsoft does it too!

- There is an application level

- There is a OS level (e.g. Windows 8)

- There is, of course, the ROM BIOS.

What is their equivalent of glibc.a/libc.c?

Fault

- ◆ Waiting on some system-allocated resource:
 - ◆ Example: memory (in a virtual memory environment):

```
int main()  
{  
    int bigArray[5000];  
  
    bigArray[4999] = 4998;  
    // Page fault when not in memory  
}
```

Abort

- ◆ OS “You can't do *that!*”
 - ◆ Writing to where you can't write to
 - ◆ Reading from where you can't read from

```
int main()  
{  
    int bigArray[5000];  
  
    bigArray[-400000] = 4998;  
    // Abort when addr not in page table  
}
```

Your turn!

Can't the compiler or linker tell beforehand whether an instruction will be a trap or abort and do something about it at compile or link time?

Yeah yeah, other than `printf()` what other system calls can I do?

- Among the most powerful things you can do are:
 - Make *Brand New* processes! (This lecture)
 - Tell your processes how to behave when they receive *signals*! (Next lecture)

`fork()`, `getpid()` and `getppid()`

- ◆ `fork()` makes a baby process that is ***an exact copy*** of a mama process
- ◆ `fork()` returns type `pid_t` (usually an `int`):
 - ◆ -1: ***“Too many processes, fool!”***
 - ◆ 0: ***“I must be the baby”***
 - ◆ (positive number): ***“I'm the mama and I got my baby's number”***
- ◆ **Question:** How does a baby know its number?
 - ◆ **Answer:** `getpid()`
- ◆ **Question:** How does a baby know its mama?
 - ◆ **Answer:** `getppid()`

fork() example 1

```
#include <stdlib.h>
#include <stdio.h>

int main ( )
{
    int x = 1;
    pid_t pid = fork();

    if (pid == 0)
        printf("Child has x = %d\n", ++x);
    else
        printf("Parent has x = %d\n", --x);

    printf("Bye from process %d with x = %d\n",
           getpid(), x);
    return(0);
}
```

fork() example 2

// How many processes result from this?

```
#include <stdlib.h>
```

```
#include <stdio.h>
```

```
int main ( )
```

```
{
```

```
    puts( "L0" );
```

```
    fork( );
```

```
    puts( "L1" );
```

```
    fork( );
```

```
    printf( "Bye from process %d\n", getpid( ) );
```

```
    return( 0 );
```

```
}
```

Your turn!

- ◆ OS's are cheap and lazy. If they can get away with letting the mama and baby process share a segment, they will.
- ◆ Which segments can mama and baby share?
 - ◆ `.text`
 - ◆ `.rodata`
 - ◆ `.data` and `.bss`
 - ◆ Heap
 - ◆ Dynamic libraries
 - ◆ Stack

Your turn, again!

- ♦ Write programs to prove that the mama and baby have distinct:
 - ♦ `.data` segments
 - ♦ `.bss` segments
 - ♦ Stacks
 - ♦ Heaps

`exit ()`, a way to end the process

- ◆ `exit (EXIT_SUCCESS)`
 - ◆ Everything went fine
 - ◆ Integer value 0
- ◆ `exit (EXIT_FAILURE)`
 - ◆ *Uhoh!* An error occurred.
 - ◆ Integer value 1
- ◆ Also doable as `return (EXIT_SUCCESS)` or `return (EXIT_FAILURE)` within `main ()`

atexit () : A way to do things after main () ends

```
#include <stdlib.h>
#include <stdio.h>

void cleanup ( )
{
    printf("Cleaning up after %d\n",getpid());
}

int main ( )
{
    atexit(cleanup);
    puts("Forking");
    fork();
    printf("Process %d's main() finished.\n",getpid());
    return(0);
}
```


exec1 () : same process, different program

- ◆ Same process running different program:

```
#include <stdlib.h>
#include <stdio.h>
#include <string.h>
#include <unistd.h>
int      main      ( )
{
    char  line[100];
    printf("What program do you want to run? ");
    fgets(line,100,stdin);
    char* cPtr = strchr(line,'\n');
    if  (cPtr != NULL)
        *cPtr = '\0';
    exec1(line,line,NULL);
    return(EXIT_SUCCESS);
}
```

Your turn!

`fork()` and `exec1()` together:

Write a program where the mama process asks for a program to run, and the baby process runs it.

“But wait! There's more . . .”

- ◆ There are several versions of `exec* ()`.
- ◆ All do the same thing with different parameters:
 - ◆ E.g. `execv ()` lets you pass a NULL-terminated array instead of listing all arguments in command:
 - ◆ Do `$ man exec1` for more details

```
#include <unistd.h>
```

```
extern char **environ;
```

```
int execl(const char *path, const char *arg, ...);  
int execlp(const char *file, const char *arg, ...);  
int execl_e(const char *path, const char *arg, ...,  
            char * const envp[]);  
int execv(const char *path, char *const argv[]);  
int execvp(const char *file, char *const argv[]);  
int execvpe(const char *file, char *const argv[],  
            char *const envp[]);
```

Almost a shell!

- ◆ We have to tell the mama process to *wait* for the baby process
- ◆ `pid_t wait(int* statusIntPtr)`
 - ◆ How fortuitous!
 - ◆ Make mama wait until some baby finishes
 - ◆ Return the process id of the baby that finishes
 - ◆ Sets pointed-to address equal to an integer that encodes `exit()` integer *and* if it exited properly
 - ◆ `WIFEXITED(statusInt)`: non-0 means “okay”, 0 means “error”
 - ◆ `WEXITSTATUS(statusInt)`: Lowest byte of integer returned by `exit()`.

Your turn again!

Revise our almost-shell program so that the
mama `wait ()`s until for baby to finish.

Waiting for the *right* baby

- `wait()` waits until some baby (the first baby) finishes
- Want to wait for a ***specific*** baby?

```
pid_t waitpid(pid_t childId, int* statusIntPtr, 0)
```
- Want ***not*** to wait ***if no baby has finished?***

```
pid_t waitpid(pid_t childId, int* statusIntPtr, WNOHANG)
```
- Want to wait for ***any old baby*** just like ***wait()***?

```
pid_t waitpid(-1, int* statusIntPtr, 0)
```

waitpid() example, pg 1

```
/* Setup */
#include <stdlib.h>
#include <stdio.h>
#include <unistd.h>
#include <sys/types.h>
#include <sys/wait.h>

const int N = 16;

int main ( )
{
    pid_t pid[N];
    int    i;
    int    status;
```

waitpid() example, pg 2

```
/* Making babies
   How many second does the FIRST baby wait?
   How many second does the LAST baby wait?
*/
for (i = 0; i < N; i++)
    if ((pid[i] = fork()) == 0)
    {
        /* Child case */
        sleep(N-i);
        return(100+i);
    }
```


waitpid() example, pg 3

```
/* waiting for babies (babysitting?) */
for (i = 0; i < N; i++)
{
    pid_t childPid = waitpid(pid[i], &status, 0);

    if (WIFEXITED(child_status))
        printf("Child %d ended with status %d\n",
               childPid, WEXITSTATUS(status));
    else
        printf("Child %d ended abnormally\n",
               childPid);
}

return(EXIT_SUCCESS);
}
```

Child neglect

1. Mama `fork()`s a baby
2. The baby does its job and finishes
3. Baby takes no CPU time (it's finished) but there is still an entry in process table (so mama can get return status, *etc.*)
4. But mama ***ignores*** finished baby



The Baby becomes a ZOMBIE!



Zombie survival guide

- ♦ A finished process with an entry in process table is “not-quite” dead. It's a zombie!
- ♦ The not-so-bad:
 - ♦ Zombies take no CPU time (they're finished)
- ♦ The pretty-bad:
 - ♦ Zombies do take System memory (in process table)
 - ♦ **One zombie?** *No big deal.*
 - ♦ **An army of zombies?** *System starts running*
vvvveerrrryyy ssslllloooooowwwllllyyy
 - ♦ Might have to reboot system (*Is this a problem with any OS you know? ;)*

zombie_maker.c, page 1

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <sys/types.h>
#include <sys/wait.h>
```

```
int main ()
```

```
{
```

```
    int child_status;
```

```
    if (fork() == 0)
```

```
    {
```

```
        printf("HC: hello from child %d\n",getpid());
```

```
        return(50);
```

```
    }
```

zombie_maker.c, page 2

```
else
{
    printf("HP: hello from parent\n"
           "In another shell say 'ps aux'\n"
           "Look for '<defunct>'.\n"
           );
    sleep(20);
    wait(&child_status);
    printf("CT: child has terminated and has given\n"
           "us %d\n",
           WEXITSTATUS(child_status)
           );
}
printf("Bye\n");
return(EXIT_SUCCESS);
}
```


Next time: *Signals!*