

CSC 374/407: Computer Systems II

Lecture 5

Joseph Phillips
De Paul University

2014 August 5

Copyright © 2011-2014 Joseph Phillips
All rights reserved

Reading

- ♦ Bryant & O'Hallaron “*Computer Systems, 2nd Ed.*”
 - Chapter 13.1-13.5: Concurrent Programming
- ♦ Hoover “*System Programming*”
 - (None)

Topics

- ◆ What threads are and why use them
- ◆ How to create threads
- ◆ Critical sections and unsafe thread programming
- ◆ Synchronization
- ◆ Application: Producer/Consumer

What threads are and why use them

Purpose: to have multiple “little processes” parallelize action of one program

Threads are like processes, but they share:

- Code segment
- Heap segment
- Data/BSS segments
- Stack segment

They differ in that:

- Have different registers (including condition codes)
- Start in different regions of stack
- Have unique thread id

How to compile for threads

At the beginning of your program:

```
. . .  
#include <thisHeader.h>  
#include <pthread.h>  
#include <thatHeader.h>  
. . .
```

When you link:

```
unix> gcc . . . -lpthread . . .
```

Why are they called “p-threads?”

- **P** = **POSIX** = **P**ortable **O**perating **S**ys **I**nterface for **UniX**
- Standardizes “Unix” across Linux, BSD, Solaris, *etc.*

How to create threads

```
int pthread_create
(pthread_t* restrict          thread,
 /* Pointer to a pthread_t object to identify
  the child thread */
const pthread_attr_t* restrict attr,
 /* Pointer to optional object for properties
  of child.  You can just say NULL. */
void *(*fncName)(void*),
 /* Name of function to run:
  void* fncName(void* ptr) */
void *restrict               arg
 /* Ptr to object that is arg to fncName() */
)
```

Return value is:

- 0: success
- anything else: ERROR!

How to wait for threads

```
int pthread_join
    (pthread_t      thread,
     /* Thread for which to wait */
     void**         valuePtr
     /* Value returned by pthread_exit()
        pointed to by valuePtr (may be NULL)
        */
    )
```

Compare threads with fork:

```
/* I use fork() to create a process */
```

```
#include <stdlib.h>
```

```
#include <stdio.h>
```

```
#include <sys/types.h>
```

```
#include <sys/wait.h>
```

```
int main()
```

```
{
```

```
    int    childsStatus;
```

```
    pid_t childId  = fork();
```

```
    char* argsText = /* whatever */
```

```
    if (childId == 0)
```

```
        execlp("doThisFile",...,argsText,.. );
```

```
    waitpid(childId,&childStatus,0);
```

```
    return(0);
```

```
}
```

```
/* I use pthread_create() */
```

```
#include <stdlib.h>
```

```
#include <stdio.h>
```

```
#include <pthread.h>
```

```
int main()
```

```
{
```

```
    pthread_t    childId;
```

```
    SendObject    arg = /*whatever*/
```

```
    ReceiveObject* childStatusPtr;
```

```
    pthread_create(&childId,NULL,  
                  doThisFnc,&arg  
                  );
```

```
    pthread_join(childId,&childStatusPtr);
```

```
    return(0);
```

```
}
```


Example program:

```
/* Compile with:
 * gcc -lpthread thread_ex1.c -o thread_ex1
 */
#include <stdlib.h>
#include <stdio.h>
#include <string.h>
#include <pthread.h>
const int N = 2;

/* thread routine */
void *thread_routine(void *vargp)
{
    int id = *(int*)vargp;
    char* ptr;
    printf("Hello from child thread %d\n", id);
    switch (id)
    {
        case 0: ptr = strdup("Hello "); break;
        case 1: ptr = strdup("there!"); break;
    }
    return(ptr);
}
```

```
int main()
{
    int i;
    char* msgPtr;
    pthread_t tid[N];

    for (i = 0; i < N; i++)
        pthread_create
            (&tid[i], NULL,
             thread_routine, (void *)&i);

    for (i = 0; i < N; i++)
    {
        pthread_join(tid[i], (void**)&msgPtr);
        puts(msgPtr);
        free(msgPtr);
    }

    return(0);
}
```

Time for you!

Can you write a threaded program that can:

1. **Prove** that all threads use the same **stack** (did we just do that?)
2. **Prove** that all threads use the same **global var data space** (for global vars and static vars inside functions)
3. **Prove** that all threads use the same **heap**.

What's wrong with this?

```
/* Compile with:
 * gcc -lpthread badcnt.c -o badcnt
 */
#include <stdlib.h>
#include <stdio.h>
#include <pthread.h>

unsigned int cnt = 0; /* shared */
unsigned int NUM_ITERS;

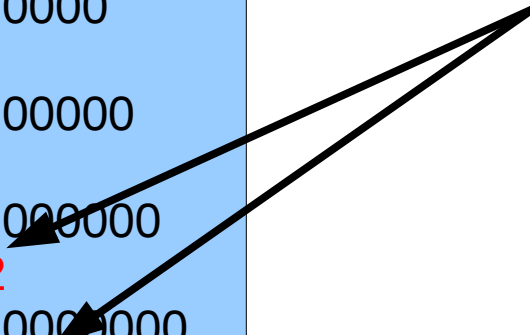
/* thread routine */
void *count(void *arg)
{
    int i;
    for (i=0; i<NUM_ITERS; i++)
        cnt++;
    return NULL;
}
```

```
int main(int argc, char* argv[])
{
    pthread_t tid1, tid2;
    if ( (argc >= 2) && isdigit(*argv[1]) )
        NUM_ITERS = atoi(argv[1]);
    else
    {
        const int LINE_SIZE = 16;
        char line[LINE_SIZE];
        do {
            printf("How many iterations? ");
            fgets(line, LINE_SIZE, stdin);
        }
        while ( !isdigit(line[0]) );
        NUM_ITERS = atoi(line);
    }
    pthread_create(&tid1, NULL, count, NULL);
    pthread_create(&tid2, NULL, count, NULL);
    pthread_join(tid1, NULL);
    pthread_join(tid2, NULL);
    printf("Should be %d is %d\n", NUM_ITERS*2, cnt);
    return(0);
}
```

Output:

```
[jphillips@localhost]$ badcnt 10  
Should be 20 is 20  
[jphillips@localhost]$ badcnt 100  
Should be 200 is 200  
[jphillips@localhost]$ badcnt 1000  
Should be 2000 is 2000  
[jphillips@localhost]$ badcnt 10000  
Should be 20000 is 20000  
[jphillips@localhost]$ badcnt 100000  
Should be 200000 is 200000  
[jphillips@localhost]$ badcnt 1000000  
Should be 2000000 is 1830352  
[jphillips@localhost]$ badcnt 10000000  
Should be 20000000 is 15214384
```

Last two
values are
WRONG!



Critical sections

(gdb) disass count

```
0x400688: push  %rbp
0x400689: mov   %rsp,%rbp
0x40068c: mov   %rdi,0xffffffffffe8(%rbp)
0x400690: movl  $0x0,0xffffffffffc(%rbp)
0x400697: jmp   0x4006ac <count+36>
0x400699: mov   2098629(%rip),%eax # 0x600c64 <cnt>
0x40069f: add   $0x1,%eax
0x4006a2: mov   %eax,2098620(%rip) # 0x600c64 <cnt>
0x4006a8: addl  $0x1,0xffffffffffc(%rbp)
0x4006ac: mov   0xffffffffffc(%rbp),%edx
0x4006af: mov   2098611(%rip),%eax #0x600c68 <NUM_ITERS>
0x4006b5: cmp   %eax,%edx
0x4006b7: jb    0x400699 <count+17>
0x4006b9: mov   $0x0,%eax
0x4006be: leaveq
0x4006bf: retq
```



These
three
must
be done
by one
thread at
a time!

Critical sections (2)

Imagine some scenarios where they are **not** done atomically

Other ways threads can step on each other's toes

1. Working with the same global variable
 - (Just covered)
2. Working with the same static var within a fnc.
 - Can *you* think of an example?
3. A function returning a pointer to a static var.
 - Can *you* think of an example?
4. Calling a thread-unsafe function.
 - Can *you* think of an example?

Functions: safe and unsafe

Reentrant functions call no shared variables

- Always thread safe

Most standard C Library functions are thread safe

- e.g. printf(), malloc()

Unsafe to safe

- asctime() (***unsafe!***)
- ctime() (***unsafe!***)
- gethostbyaddr() (***unsafe!***)
- inet_ntoa(***unsafe!***)
- localtime() (***unsafe!***)
- rand() (***unsafe!***)

asctime_r() (***safe!***)
ctime_r() (***safe!***)
gethostbyaddr_r()(***safe!***)
NO SAFE VERSION!
localtime_r() (***safe!***)
rand_r() (***safe!***)

The semaphore solution

Classic solution: Dijkstra's P and V operations on semaphores.

- ***semaphore***: non-negative integer synchronization variable.
 - `sem_wait(s): [while(s==0) wait(); s--;]`
 - Originally named P(), Dutch for "Proberen" (test)
 - `sem_post(s): [s++;]`
 - Originally named V(), Dutch for "Verhogen" (increment)
- OS guarantees that operations between brackets `[]` are executed indivisibly.
 - Only one P or V operation at a time can modify s.
 - When while loop in P terminates, only that P can decrement s.

Semaphore invariant: $(s \geq 0)$

(1) POSIX Semaphores

What to include:

- `#include <semaphore.h>`

Types and functions:

- `sem_t semaphore;`
- `sem_init(sem_t* semPtr, int flag, int value)`
 - Initialize pointed-to semaphore, with `value`, if `flag == 1` then semaphore can be forked
- `sem_destroy(sem_t* semPtr)`
 - Destroy pointed-to semaphore. If it's negative then block.

POSIX Semaphores, cont'd

- `sem_wait(sem_t* semPtr)`
 - Decrement pointed-to semaphore. If it's negative then block.
- `sem_post(sem_t* semPtr)`
 - Increment pointed-to semaphore. Wake one blocked process if any.
- `sem_getvalue(sem_t* semPtr, int* valuePtr)`
 - Get value of pointed to semaphore.

POSIX semaphore solution

```
sem_t      sem;
```

```
/* P operation on semaphore sem */
```

```
void P(sem_t *semPtr) {  
    if (sem_wait(semPtr)) {  
        fputs("sem_wait error\n",stderr);  
        exit(-1);  
    }  
}
```

```
/* V operation on semaphore sem */
```

```
void V(sem_t *semPtr) {  
    if (sem_post(semPtr)) {  
        fputs("sem_post error\n",stderr);  
        exit(-1);  
    }  
}
```

```
/* thread routine */
```

```
void *count(void *arg) {  
    int i;  
    for (i=0; i<NUM_ITERS; i++) {  
        P(&sem);  
        cnt++;  
        V(&sem);  
    }  
    return NULL;  
}
```

```
/* in main */
```

```
...
```

```
/* initialize sem to 1 */
```

```
if (sem_init(&sem, 0, 1) < 0) {  
    fputs("sem_init error\n",stderr);  
    exit(1);  
}
```

(2) *pthread_mutex* solution

- `pthread_mutex_t lock`
- `pthread_mutex_init(addressOfLock, NULL)`
 - Makes a lock and initializes it to default values.
- `pthread_mutex_destroy(addressOfLock)`
 - Makes a lock and initializes it to default values.
- `pthread_mutex_lock(address of lock)`
 - Blocks thread until lock obtained, then obtains lock and blocks other threads until lock released.
- `pthread_mutex_unlock(address of lock)`
 - Releases lock allowing other threads to obtain it.

pthread_mutex solution

```
pthread_mutex_t cnt_lock;

/* thread routine */
void *count(void *arg) {
    pthread_mutex_t* mutexLockPtr;
    mutexLockPtr = (pthread_mutex_t*)arg;

    for (int i=0; i<NUM_ITERS; i++) {
        pthread_mutex_lock(mutexLockPtr);
        cnt++;
        pthread_mutex_unlock(mutexLockPtr);
    }
    return NULL;
}

/* in main() */
...
pthread_mutex_init(&cnt_lock,NULL);
```

```
[jphillips@localhost]$ bettercnt 10
Should be 20 is 20
[jphillips@localhost]$ bettercnt 100
Should be 200 is 200
[jphillips@localhost]$ bettercnt 1000
Should be 2000 is 2000
[jphillips@localhost]$ bettercnt 10000
Should be 20000 is 20000
[jphillips@localhost]$ bettercnt 100000
Should be 200000 is 200000
[jphillips@localhost]$ bettercnt 1000000
Should be 2000000 is 2000000
[jphillips@localhost]$ bettercnt 10000000
Should be 20000000 is 20000000
[jphillips@localhost]$ bettercnt 100000000
Should be 200000000 is 200000000
```

Producer-Consumer

One or more threads produce something

- Place in buffer

One or more threads consume something

- Retrieve from buffer

Critical section

- Access of pointers/indices for buffer

Problem!

- A producer may gain access to a full buffer
- A consumer may gain access to an empty buffer

Oh no! An example in C++!

- A buffer is an object.
- Make the object thread-safe.
- Lessons carry over to Java, C#, etc.

Unsafe_Buffer.h

```
class Buffer
{
    enum { SIZE = 16 };

    int array_[SIZE];
    int inIndex_;
    int outIndex_;
    int numItems_;

public :

    Buffer ()
    {
        inIndex_ = outIndex_
            = numItems_ = 0;
    }

    ~Buffer ()
    {
    }

    int getNumItems () const
    { return(numItems_); }

    void putIn (int i)
    {
        while (getNumItems() >= SIZE)
        {
            printf("Full! Waiting!\n");
            usleep(10);
        }

        array_[inIndex_] = i;
        countArray[array_[inIndex_]]++;
        usleep(10 + rand() % 10);
        inIndex_++;
        numItems_++;
        if (inIndex_ >= SIZE)
            inIndex_ = 0;
    }
}
```


Unsafe_Buffer.h

```
int    pullOut ()
{
    while  (getNumItems() <= 0)
    {
        printf("Empty!   Waiting!\n");
        usleep(10);
    }

    countArray[array_[outIndex_]]--;
    int toReturn      = array_[outIndex_];
    usleep(10 + rand() % 10);

    outIndex_++;
    numItems_--;
    if  (outIndex_ >= SIZE)
        outIndex_ = 0;

    return(toReturn);
}
};
```

producerConsumer.cpp

```
#include          <stdlib.h>
#include          <stdio.h>
#include          <unistd.h>
#include          <pthread.h>
int*
countArray;
#include          "Unsafe_Buffer.h"

const int        NUM_INTEGERS_TO_BUFFER  = 0x1000;

void*    stuffIntegersIn (void*    vPtr)
{
    for  (int i = 0;   i < NUM_INTEGERS_TO_BUFFER;   i++)
        ((Buffer*)vPtr)->putIn(i);

    return(NULL);
}
```

producerConsumer.cpp

```
void*    pullIntegersOut (void*   vPtr)
{
    for  (int i = 0;   i < NUM_INTEGERS_TO_BUFFER;   i++)
    {
        int j = ((Buffer*)vPtr)->pullOut();

        printf("Trial %d got %d.\n",i,j);
        fflush(stdout);
    }

    return(NULL);
}
```

producerConsumer.cpp

```
int      main      ( )
{
    pthread_t      producer0;
    pthread_t      producer1;
    pthread_t      consumer0;
    pthread_t      consumer1;
    Buffer          buffer;
    countArray      =
        (int*)calloc(NUM_INTEGERS_TO_BUFFER, sizeof(int));

    pthread_create(&producer0, NULL, stuffIntegersIn, &buffer);
    pthread_create(&producer1, NULL, stuffIntegersIn, &buffer);
    pthread_create(&consumer0, NULL, pullIntegersOut, &buffer);
    pthread_create(&consumer1, NULL, pullIntegersOut, &buffer);
}
```

producerConsumer.cpp

```
pthread_join(producer1,NULL);
pthread_join(producer0,NULL);
pthread_join(consumer1,NULL);
pthread_join(consumer0,NULL);

for (int i = 0; i < NUM_INTEGERS_TO_BUFFER; i++)
    if (countArray[i] < 0)
        printf("%d was gotten too many times!\n",i);
    else
        if (countArray[i] > 0)
            printf("%d was put too many times!\n",i);

return(EXIT_SUCCESS);
}
```

Your turn!

(1) Run it as-is. *Any problems?*

Your turn!

(2) Make it thread-safe.

```
void myMethod ( )  
{  
    pthread_mutex_lock(&lock);  
    doCriticalSection();  
    pthread_mutex_unlock(&lock);  
}
```

NOTE:

- 1) Each object gets its **own** lock.
- 2) All threads accessing the same object use **same** lock

Any problems still?

Solution: *pthread_cond*!

- `pthread_cond_t cond`
- `pthread_cond_init(addressOfCondition, NULL)`
 - Makes a condition and initializes it to default values.
- `pthread_cond_destroy(addressOfCondition)`
 - Destroys condition.
- `pthread_cond_wait(addrOfCondition, addrOfLock)`
 - Blocks thread until lock released and condition signaled. Then obtains lock again.
- `pthread_cond_signal(addressOfCondition)`
 - Signal *one waiting thread* that condition is met.
- `pthread_cond_broadcast(addressOfCondition)`
 - Signal *all waiting threads* that condition is met.

The proper solution

```
void myMethod ( )
{
    pthread_mutex_lock(&lock);
    while ( !this->isReady() )
        pthread_cond_wait(&wCond,&lock);
    doCriticalSection();
    pthread_cond_signal(&sCond);
    pthread_mutex_unlock(&lock);
}
```

NOTE:

- (1) Method **isReady()** determines if this object is ready for thread to do **myMethod()**.
- (2) Conditions **wCond** and **sCond** may or may not be same.

We have mutual exclusion, are guaranteed to be safe?

Sorry but No!

- Race conditions
 - When result is dependent on order of processes or threads
- Deadlock
 - Thread/process A has lock 1 and is waiting for lock 2
 - Thread/process B has lock 2 and is waiting for lock 1

Next time: *Memory!*